

I Introduction

I.A General Information

- Addresses are 32-bit.
- Little-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
 - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be **0b0**

I.B Registers

There are sixteen general-purpose registers (all of which are 32-bit): **r0**, **r1**, **r2**, ..., **r11**, **r12**, **lr** (link register, the return address of **bl** and **jl**), **fp** (frame pointer), and **sp** (stack pointer).

The program counter, **pc**, is 32 bits long, as addresses are 32-bit.

Here are the special-purpose registers:

- **flags**: arithmetic/logic **FLAGS**; (reg encoding: **0x0**)
- **ids**: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: **0x1**)
- **ira**: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: **0x2**)
- **ie**: IRQ Enable flag: flag indicating whether IRQs are disabled (**0x0**) or enabled (**0x1**); note that this flag starts with a value of **0x0**; (reg encoding: **0x3**)
- **ity**: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (**0x0**) or a **swi** (**0x1**) (reg encoding: **0x4**)
- **sty**: Software interrupt TYpe: **swi**'s argument. For **swi rA, #simm**, this is the value **rA + simm** (reg encoding: **0x5**)
- **hi**:
 - Uses:
 - high 32 bits of **lumul** and **lsmul** multiply result
 - high 32 bits of 64-bit **divmod** modulo result
 - (reg encoding: **0x6**)
- **lo**:
 - Uses:
 - low 32 bits of **lumul** and **lsmul** multiply result
 - low 32 bits of 64-bit **divmod** modulo result
 - 32-bit **divmod** modulo result
 - (reg encoding: **0x7**)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of **flags**:

- Zero (Z): (**flags** bit 0)
- Carry (C): (**flags** bit 1)
- oVerflow (V): (**flags** bit 2)
- Negative (N): (**flags** bit 3)
- Note: All other bit positions of **flags** are reserved.

II Instruction Set

II.A Instruction Group 0: **pre**, **lpre**, and **atomics**

For **pre**, the following encoding is used, with each character representing one bit:

0000 iiii iiii iiii, where

- **i** is a 12-bit constant.

For `lpre`, the following encoding is used, with each character representing one bit:

0001 0iii iiiii iiiii iiiii iiiii iiiii, where

- `i` is a 27-bit constant.

For `cmpxchg` and `xchg`, the following encoding is used, with each character representing one bit:

0001 100l bbbb aaaa, where

- `l` encodes whether to lock (1) or not lock (0)
- `b` encodes register `rB`
- `a` encodes register `rA`

This instruction acts as `cmpxchg` if `index` is in effect, using `indexRegA` as the **expected** (or `rC`) value, but otherwise this instruction acts as `xchg`. More on `index` later in this document.

The assembly syntax of these instructions is as follows:

- For `l=0` (without lock):
 - `cmpxchg [rA], rC, rB`
 - `xchg rA, rB`
- For `l=1` (with lock):
 - `cmpxchg.l [rA], rC, rB`
 - `xchg.l [rA], rB`

`cmpxchg/cmpxchg.l` sets the Z flag to 1 upon success and 0 upon failure.

II.A.1 `pre` and `lpre`

`pre` and `lpre` are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use `pre` or `lpre` as instructions. Instead, it is expected that the assembler or linker will be the one to insert `pre` or `lpre` as needed if an immediate is too large for a particular instruction.

For non-branch instructions:

- If `pre` is used, the immediate field of the `pre` instruction will form bits [16:5] of the immediate of the next non-`index` instruction. The 17-bit immediate will then be sign-extended to 32 bits. Sign-extension will be performed on the 17-bit immediate even if, had there been no `pre`, the 5-bit immediate would have been zero-extended.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits [31:5] of the immediate of the next non-`index` instruction.

For branch instructions (group 3):

- if `pre` is used, the immediate field of the `pre` instruction will form bits [20:9] of the immediate of the next non-`index` instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits [31:9] of the immediate of the next non-`index` instruction.

II.A.2 Handling of `pre`, `lpre`, and `index`

When a `pre` or `lpre` instruction is found, `pre` or `lpre` will be considered to be “in effect”. This condition lasts for one or two instructions after the `pre` or `lpre` instruction, depending on whether or not `index` was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + indexRegSum` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre/lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, IRQs will not be serviced.

Pseudo code for handling the how pre, lpre, and index are to be handled regarding whether or not they are “in effect” is as follows:

```
function handlePreLpreIndex(input in, output out) {
  when (!out.instr.isLpre()) {
    when (in.instr.isPre()) {
      when (
        in.state.pre.haveIt
        || in.state.lpre.haveIt
      ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
        out.instr := NOP
        out.state.canServiceIrq := True
      } otherwise {
        out.state.pre.haveIt := True
        out.state.pre.data := in.instr[11:0]
        out.instr := in.instr
        out.state.canServiceIrq := False
      }
    } elseif (in.instr.isLpre()) {
      when (
        in.state.pre.haveIt
        || in.state.lpre.haveIt
      ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
        out.instr := NOP
        out.state.canServiceIrq := True
      } otherwise {
        out.state.lpre.haveIt := True
        out.instr := in.instr
        out.state.canServiceIrq := False
        out.state.lpre.data[26:16] := in.instr[10:0]
      }
    } elseif (in.instr.isIndex()) {
      when (
        in.state.index.haveIt
      ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
        out.instr := NOP
        out.state.canServiceIrq := True
      } otherwise {
        out.state.index.haveIt := True
        out.instr := in.instr
        out.state.canServiceIrq := False
      }
    } otherwise {
      out.state.canServiceIrq := !(
        in.state.pre.haveIt
        || in.state.lpre.haveIt
        || in.state.index.haveIt
      )
      out.instr := instr
    }
  }
}
```

```

        // Whenever we see an instruction other than `pre`, `lpre`, or `index`,
        // that means those instructions stop being "in effect".
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
    }
} otherwise { // when (out.instr.isLpre())
    out.state.lpre.data[15:0] := in.instr
}
}

```

II.B Instruction Group 1

The following encoding is used, with each character representing one bit:

001i iiii oooo aaaa, where

- i is a 5-bit sign-extended or zero-extended immediate, and is denoted `simm` when sign-extended or `imm` when zero-extended. Also, `simm` or `imm` can be expanded with `pre` or `lpre`.
- a encodes register `rA`
- o is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: `add rA, #simm`
- Opcode 0x1: `add rA, pc, #simm`
 - Effect: `rA <= pc + simm + 2;`
- Opcode 0x2: `add rA, sp, #simm`
- Opcode 0x3: `add rA, fp, #simm`
- Opcode 0x4: `cmp rA, #simm`
 - Effect: Compare `rA` to `simm`.
- Opcode 0x5: `cpy rA, #simm`
 - Effect: Copy an immediate value into `rA`
- Opcode 0x6: `lsl rA, #imm`
 - Effect: Logical shift left
- Opcode 0x7: `lsr rA, #imm`
 - Effect: Logical shift right
- Opcode 0x8: `asr rA, #imm`
 - Effect: Arithmetic shift right
- Opcode 0x9: `and rA, #simm`
 - Effect: Bitwise AND
- Opcode 0xa: `orr rA, #simm`
 - Effect: Bitwise OR
- Opcode 0xb: `xor rA, #simm`
 - Effect: Bitwise XOR
- Opcode 0xc: `ze rA, #imm`
 - Effect: Set `rA[31:imm]` to zero.
- Opcode 0xd: `se rA, #imm`
 - Effect: Set each bit of `rA[31:imm]` to the bit `rA[imm]`.
- Opcode 0xe: `swi rA, #simm`
 - Effect: Call software interrupt number `rA + simm`.
- Opcode 0xf: `swi #imm`
 - Effect: Call software interrupt number `imm`.

II.C Instruction Group 2

The following encoding is used, with each character representing one bit:

010f oooo bbbb aaaa, where

- o is the opcode
- b encodes register `rB`

- **a** encodes register **rA**
- **f** is encoded as **0** if this instruction cannot affect flags and encoded **1** if this instruction is permitted to affect flags. Note that **cmp** is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.

- Opcode **0x0**: **add rA, rB**
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
- Opcode **0x1**: **sub rA, rB**
 - Mnemonic for when flags not affected: **sub**
 - Mnemonic for when flags affected: **sub.f**
- Opcode **0x2**: **add rA, sp, rB**
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
- Opcode **0x3**: **add rA, fp, rB**
 - Mnemonic for when flags not affected: **add**
 - Mnemonic for when flags affected: **add.f**
- Opcode **0x4**: **cmp rA, rB**
 - Effect: Compare **rA** to **rB**. **cmp** is *always* able to affect flags, independent of the encoded **f** bit of the instruction.
- Opcode **0x5**: **cpy rA, rB**
 - Mnemonic for when flags not affected: **cpy**
 - Mnemonic for when flags affected: **cpy.f**
 - Effect: Copy **rB** into **rA**
- Opcode **0x6**: **lsl rA, rB**
 - Mnemonic for when flags not affected: **lsl**
 - Mnemonic for when flags affected: **lsl.f**
 - Effect: Logical shift left
- Opcode **0x7**: **lsr rA, rB**
 - Mnemonic for when flags not affected: **lsr**
 - Mnemonic for when flags affected: **lsr.f**
 - Effect: Logical shift right
- Opcode **0x8**: **asr rA, rB**
 - Mnemonic for when flags not affected: **asr**
 - Mnemonic for when flags affected: **asr.f**
 - Effect: Arithmetic shift right
- Opcode **0x9**: **and rA, rB**
 - Mnemonic for when flags not affected: **and**
 - Mnemonic for when flags affected: **and.f**
 - Effect: Bitwise AND
- Opcode **0xa**: **orr rA, rB**
 - Mnemonic for when flags not affected: **orr**
 - Mnemonic for when flags affected: **orr.f**
 - Effect: Bitwise OR
- Opcode **0xb**: **xor rA, rB**
 - Mnemonic for when flags not affected: **xor**
 - Mnemonic for when flags affected: **xor.f**
 - Effect: Bitwise XOR
- Opcode **0xc**: **adc rA, rB**
 - Mnemonic for when flags not affected: **adc**
 - Mnemonic for when flags affected: **adc.f**
 - Effect: Add with Carry, using the formula **rA + rB + flags.C** to compute the value that will be written into **rA**.

- Opcode 0xd: `sbc rA, rB`
 - Mnemonic for when flags not affected: `sbc`
 - Mnemonic for when flags affected: `sbc.f`
 - Effect: Subtract with Borrow, using the formula $rA + (\sim rB) + \text{flags.C}$ to compute the value that will be written into `rA`.
- Opcode 0xe: `cmpbc rA, rB`
 - Effect: Compare `rA` to `rB`, but with carry-in and a different effect for setting the `flags.Z`. `cmpbc` is *always* able to affect `flags`, independent of the encoded `f` bit of the instruction.
 - Note: this instruction acts much like `sbc rA, rB`, but without storing the subtraction's result into `rA`. However, this instruction sets the Z flag to $\text{prev}(\text{flags.Z}) \text{ AND } ((rA + (\sim rB) + \text{flags.C}) == 0)$

II.D Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:

011i iiii iiii oooo, where

- `i` is a 9-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `o` is the opcode

Here is a list of instructions from this encoding group.

- Opcode 0x0: `bl simm`
 - Name: Branch and Link
 - Description: Relative call
 - Effect: $lr \leq pc + 2$; $pc \leq pc + simm + 2$;
- Opcode 0x1: `bra simm`
 - Name: BRanch Always
 - Description: Unconditional relative branch
 - Effect: $pc \leq pc + simm + 2$;
- Opcode 0x2: `beq simm`
 - Name: Branch if EQual
 - Effect: if (`flags.Z`) $pc \leq pc + simm + 2$;
- Opcode 0x3: `bne simm`
 - Name: Branch if Not Equal
 - Effect: if ($\neg \text{flags.Z}$) $pc \leq pc + simm + 2$;
- Opcode 0x4: `bmi simm`
 - Name: Branch if MInus
 - Effect: if (`flags.N`) $pc \leq pc + simm + 2$;
- Opcode 0x5: `bpl simm`
 - Name: Branch if PPlus
 - Effect: if ($\neg \text{flags.N}$) $pc \leq pc + simm + 2$;
- Opcode 0x6: `bvs simm`
 - Name: Branch if oVerflow Set
 - Effect: if (`flags.V`) $pc \leq pc + simm + 2$;
- Opcode 0x7: `bvc simm`
 - Name: Branch if oVerflow Clear
 - Effect: if ($\neg \text{flags.V}$) $pc \leq pc + simm + 2$;
- Opcode 0x8: `bgeu simm`
 - Name: Branch if Greater than or Equal Unsigned
 - Effect: if (`flags.C`) $pc \leq pc + simm + 2$;
- Opcode 0x9: `bltu simm`
 - Name: Branch if Less Than Unsigned
 - Effect: if ($\neg \text{flags.C}$) $pc \leq pc + simm + 2$;
- Opcode 0xa: `bgtu simm`
 - Name: Branch if Greater Than Unsigned
 - Effect: if (`flags.C AND !flags.Z`) $pc \leq pc + simm + 2$;

- Opcode 0xb: bleu *sim*m
 - Name: Branch if Less than or Equal Unsigned
 - Effect: if (!flags.C OR flags.Z) *pc* <= *pc* + *sim*m + 2;
- Opcode 0xc: bges *sim*m
 - Name: Branch if Greater than or Equal Signed
 - Effect: if (!(flags.N XOR flags.V)) *pc* <= *pc* + *sim*m + 2;
- Opcode 0xd: blts *sim*m
 - Name: Branch if Less Than Signed
 - Effect: if (flags.N XOR flags.V) *pc* <= *pc* + *sim*m + 2;
- Opcode 0xe: bgts *sim*m
 - Name: Branch if Greater Than Signed
 - Effect:


```
if (!(flags.N XOR flags.V) AND !flags.Z)
    pc <= pc + simm + 2;
```
- Opcode 0xf: bles *sim*m
 - Name: Branch if Less than or Equal Signed
 - Effect:


```
if ((flags.N XOR flags.V) OR flags.Z)
    pc <= pc + simm + 2;
```

II.E Instruction Group 4

The following encoding is used, with each character representing one bit:

1000 0000 *bbbb* *aaaa*, where

- *o* is the opcode
- *b* encodes register *rB* or register *sC*, where *sC* is one of the special-purpose registers
- *a* encodes register *rA* or register *sA*, where *sA* is one of the special-purpose registers

Here is a list of instructions from this encoding group.

- Opcode 0x0: jl *rA*
 - Effect: *lr* <= *pc* + 2; *pc* <= *rA*;
- Opcode 0x1: jmp *rA*
 - Effect: *pc* <= *rA*;
- Opcode 0x2: jmp *ira*
 - Effect: *pc* <= *ira*;
- Opcode 0x3: reti
 - Effect: enables IRQs (by copying 0x1 into *ie*) and performs *pc* <= *ira*;
- Opcode 0x4: ei
 - Effect: copy 1 into *ie*.
- Opcode 0x5: di
 - Effect: copy 0 into *ie*.
- Opcode 0x6: push *rA*, *rB*
 - Effect: pushes *rA* onto the stack, using *rB* as the stack pointer, post-decrementing *rB*.
 - This instruction does nothing when *rA* is the same register as *rB*.
 - Note: As a pseudo instruction, omitting “, *rB*” will automatically select *sp* as the particular stack pointer.
- Opcode 0x7: push *sA*, *rB*
 - Effect: pushes *sA* onto the stack, using *rB* as the stack pointer, post-decrementing *rB*.
 - Note that *sA* is considered to be 32-bit for the purpose of the store to memory and decrementing *rB*, even if *sA* is *flags* or *ie*.
 - Note: As a pseudo instruction, omitting “, *rB*” will automatically select *sp* as the particular stack pointer.
- Opcode 0x8: pop *rA*, *rB*
 - Effect: pops *rA* off the stack, using *rB* as the stack pointer, pre-incrementing *rB*.

- This instruction does nothing when `rA` is the same register as `rB`.
- This instruction does nothing when all of the following conditions are met:
 - `rA`'s register encoding's evenness/oddness is the same as `rB`'s register encoding's evenness/oddness.
 - `rB` is **not** `sp`
- Note: As a pseudo instruction, omitting “, `rB`” will automatically select `sp` as the particular stack pointer.
- Opcode `0x9`: `pop sA, rB`
 - Effect: pops `sA` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
 - Note that `sA` is considered to be 32-bit for the purpose of the load from memory and incrementing `rB`, even if `sA` is `flags` or `ie`.
 - Note: As a pseudo instruction, omitting “, `rB`” will automatically select `sp` as the particular stack pointer.
- Opcode `0xa`: `pop pc, rB`
 - Effect: pops `pc` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
 - Note: As a pseudo instruction, omitting “, `rB`” will automatically select `sp` as the particular stack pointer.
- Opcode `0xb`: `mul rA, rB`
 - Effect: `rA <= rA * rB`;
- Opcode `0xc`: `udivmod rA, rB`
 - Effect:
 - `rA <= u32(rA) / u32(rB)`;
 - `lo <= u32(rA) % u32(rB)`;
- Opcode `0xd`: `sdivmod rA, rB`
 - Effect:
 - `rA <= s32(rA) / s32(rB)`;
 - `lo <= s32(rA) % s32(rB)`;
- Opcode `0xe`: `lumul rA, rB`
 - Effect: This instruction multiplies `rA` by `rB`, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{hi, lo}`
- Opcode `0xf`: `lsmul rA, rB`
 - Effect: This instruction multiplies `rA` by `rB`, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in `concat{hi, lo}`
- Opcode `0x10`: `udivmod64 rA, rB`
 - Effect:
 - performs a 64-bit by 64-bit unsigned division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`,
 - ...and modulo results are stored in `concat{hi, lo}`.
 - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
- Opcode `0x11`: `sdivmod64 rA, rB`
 - Effect:
 - performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`,
 - ...and modulo results are stored in `concat{hi, lo}`.
 - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
- Opcode `0x12`: `ldub rA, [rB]`
 - Effect: Load an 8-bit value from memory at the the effective address computed as either
 - `rB` (only if `index` is **not** in effect)
 - or `<indexRegSum>` (only if `index` is in effect),
 then zero-extend the 8-bit value to 32 bits, then write the zero-extended 32-bit value into `rA`.

- ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `ldub rA, [rB, rC]`
- ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `ldub rA, [rB, #simm]`
- Opcode 0x13: `ldsb rA, [rB]`
 - ▶ Effect: Load an 8-bit value from memory the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
then sign-extend the 8-bit value to 32 bits, then write the sign-extended 32-bit value into `rA`.
 - ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `ldsb rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `ldsb rA, [rB, #simm]`
- Opcode 0x14: `lduh rA, [rB]`
 - ▶ Effect: Load a 16-bit value from memory at the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
then zero-extend the 16-bit value to 32 bits, then write the zero-extended 32-bit value into `rA`.
 - ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `lduh rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `lduh rA, [rB, #simm]`
- Opcode 0x15: `ldsh rA, [rB]`
 - ▶ Effect: Load a 16-bit value from memory at the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
then sign-extends the 16-bit value to 32 bits, then writes the sign-extended 32-bit value into `rA`.
 - ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `ldsh rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `ldsh rA, [rB, #simm]`
- Opcode 0x16: `ldr rA, [rB]`
 - ▶ Effect: Store `rA` to memory at the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
 - ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `ldr rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `ldr rA, [rB, #simm]`
- Opcode 0x18: `stb rA, [rB]`
 - ▶ Effect: Store `rA[7:0]` to memory at the the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
 - ▶ Shorthand for having the assembler insert an index `rB`, `rC` instruction before this one: `stb rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB`, `#simm` instruction before this one: `stb rA, [rB, #simm]`
- Opcode 0x19: `sth rA, [rB]`
 - ▶ Effect: Store `rA[15:0]` to memory at the the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),

- ▶ Shorthand for having the assembler insert an index `rB, rC` instruction before this one: `sth rA, [rB, rC]`
- ▶ Shorthand for having the assembler insert an index `rB, #simm` instruction before this one: `sth rA, [rB, #simm]`
- Opcode `0x1a`: `str rA, [rB]`
 - ▶ Effect: Store `rA` to memory at the the effective address computed as either
 - `rB` (only if index is **not** in effect)
 - or `<indexRegSum>` (only if index is in effect),
 - ▶ Shorthand for having the assembler insert an index `rB, rC` instruction before this one: `str rA, [rB, rC]`
 - ▶ Shorthand for having the assembler insert an index `rB, #simm` instruction before this one: `str rA, [rB, #simm]`
- Opcode `0x1c`: `cpy rA, sB`
 - ▶ Effect: `rA <= sB;`
- Opcode `0x1d`: `cpy sA, rB`
 - ▶ Effect: `sA <= rB;`
- Opcode `0x1e`: `cpy sA, sB`
 - ▶ Effect: `sA <= sB;`

II.F Instruction Group 5, Subgroup **0b0**: index `rA, rB`

The following encoding is used, with each character representing one bit:

`1010 000r bbbb aaaa`, where

- `r` is a reserved bit value, which should be zero until the 64-bit version of this spec is defined.
- `b` encodes register `rB`
- `a` encodes register `rA`

The one instruction from this encoding group is index `rA, rB`

- Effect: Performs the following:
 - ▶ `indexRegSum <= indexRaSimm.haveIt ? indexRegSum + rB : rA + rB;`
 - ▶ `indexRegA <= rA;`
 - ▶ `indexRegB <= rB;`

and stores that `indexRaRb` is in effect.

- Note: If this prefix occurs when `indexRaSimm` is already in effect, then the overall instruction becomes a NOP, and `indexRaRb`, `indexRaSimm`, `pre`, and `lpre` will all no longer be in effect.

II.G Instruction Group 5, Subgroup **0b1**: index `rA, #simm`

The following encoding is used, with each character representing one bit:

`1011 riii iiii aaaa`, where

- `r` is a reserved bit value, which should be encoded as zero until the 64-bit version of this spec is defined.
- `i` is a 7-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `a` encodes register `rA`

The one instruction from this encoding group is index `rA, #simm`

- Effect: Performs the following:
 - ▶ `indexRegSum <= indexRaRb.haveIt ? indexRegSum + simm : rA + simm;`
 - ▶ `indexRegA <= rA;`
 - ▶ `indexRegB <= simm;`

and stores that `indexRaSimm` is in effect.

- Note: If this prefix occurs when `indexRaSimm` is already in effect, then the overall instruction becomes a NOP, and `indexRaRb`, `indexRaSimm`, `pre`, and `lpre` will all no longer be in effect.

II.H Instruction Group 7, Subgroup 0b00: Extra 8-bit and 16-bit Ops

The following encoding is used, with each character representing one bit:

1110 0w0o bbbb aaaa, where

- w is the operation width
 - When 0b0: 8-bit operation
 - When 0b1: 16-bit operation
- o is the opcode
- b encodes register rB
- a encodes register rA

Here is a list of instructions from this encoding group.

- Opcode 0x0:
 - w value 0b0: cmpb rA, rB
 - Effect: Compare rA[7:0] to rB[7:0]
 - w value 0b1: cmph rA, rB
 - Effect: Compare rA[15:0] to rB[15:0]
- Opcode 0x1:
 - w value 0b0: lsrB rA, rB
 - Effect: Logical shift right rA[7:0] by rB
 - w value 0b1: lsrh rA, rB
 - Effect: Logical shift right rA[15:0] by rB
- Opcode 0x2:
 - w value 0b0: asrB rA, rB
 - Effect: Arithmetic shift right rA[7:0] by rB
 - w value 0b1: asrh rA, rB
 - Effect: Arithmetic shift right rA[15:0] by rB

II.I Instruction Group 7, Subgroup 0b010: Extra load/store instructions

The following encoding is used, with each character representing one bit:

1110 100o bbbb aaaa, where

- o is the opcode
- b encodes register rB or sB
- a encodes register sA

Here is a list of instructions from this encoding group.

- Opcode 0x0: ldr sA, [rB]
 - Effect: 32-bit load of sA from memory at address held in rB.
- Opcode 0x1: ldr sA, [sB]
 - Effect: 32-bit load of sA from memory at address held in sB.
- Opcode 0x2: str sA, [rB]
 - Effect: 32-bit store of sA to memory at address held in rB.
- Opcode 0x3: str sA, [sB]
 - Effect: 32-bit store of sA to memory at address held in sB.

II.J Instruction Group 7, Subgroup 0b0110: Icache Reload Instruction

The following encoding is used, with each character representing one bit:

1110 110i iiii aaaa, where

- i is a 5-bit sign-extended immediate, which can be expanded by pre or lpre, and is denoted `simm`.
- a encodes register rA

The one instruction from this encoding group is `icreload [rA, #simm]`. This instruction forcibly reloads the icache line at an effective address computed as `rA + simm` (only if `index` is **not** in effect) or `<indexRegSum> + simm` (only if `index` is in effect), using the sign-extended form of `simm`. As a side note, this instruction will attempt to read the data from dcache, and will only read from memory if dcache doesn't have that data. Icache lines and dcache lines are the same length, 64 bytes.

The following pseudo instructions exist within the assembler for this encoding group:

- `icreload [rA]`
 - This pseudo instruction assembles to the following:
 - `icreload [rA, #0x0]`
- `icreload [rA, rC]`
 - This pseudo instruction assembles to the following:
 - `index rA, rC`
 - `icreload [r0, #0x0]`
- `icreload [rA, rC, #simm]`
 - This pseudo instruction assembles to the following:
 - `index rA, rC`
 - `icreload [r0, #simm]`

II.K Instruction Group 7, Subgroup `0b01110`: Icache Flush Instruction

The following encoding is used, with each character representing one bit:

`1110 1110 0000 0000`

The one instruction from this encoding group is `icflush`, which invalidates the instruction cache entirely.