I Introduction

I.A General Information
- Addresses are 32-bit.
- Little-endian byte ordering is used.
- Bytes are octets (8 bits).
- Instructions must be aligned to 16 bits, so jump and branch targets must also be aligned to 16 bits.
  - Branch offsets encoded into instructions must be 16-bit aligned, or in other words bit 0 of the branch offset must be `0b0`

I.B Registers
There are sixteen general-purpose registers (all of which are 32-bit): `r0`, `r1`, `r2`, ..., `r11`, `r12`, `lr` (link register, the return address of `bl` and `jl`), `fp` (frame pointer), and `sp` (stack pointer).

The program counter, `pc`, is 32 bits long, as addresses are 32-bit.

Here are the special-purpose registers:
- `flags`: arithmetic/logic FLAGS; (reg encoding: `0x0`)
- `ids`: Interrupt DeStination: the address to jump to upon an interrupt being serviced (also known as the interrupt vector); (reg encoding: `0x1`)
- `ira`: Interrupt Return Address: the address that was jumped from upon an interrupt being serviced; (reg encoding: `0x2`)
- `ie`: IRQ Enable flag: flag indicating whether IRQs are disabled (`0x0`) or enabled (`0x1`); note that this flag starts with a value of `0x0`; (reg encoding: `0x3`)
- `ity`: Interrupt TYpe: flag indicating whether the most recently taken interrupt is an IRQ (`0x0`) or a `swi` (`0x1`) (reg encoding: `0x4`)
- `sty`: Software interrupt TYpe: `swi`'s argument. For `swi rA, #simm`, this is the value `rA + simm` (reg encoding: `0x5`)
- Note: All other encodings for special-purpose registers are reserved.

Here are the bits of `flags`:
- Zero (Z): (`flags` bit 0)
- Carry (C): (`flags` bit 1)
- oVerflow (V): (`flags` bit 2)
- Negative (N): (`flags` bit 3)
- Note: All other bit positions of `flags` are reserved.

II Instruction Set

II.A Instruction Group 0: `pre`, `lpre`, and atomics
For `pre`, the following encoding is used, with each character representing one bit:
`0000 iiii iiii iiii`, where
- `i` is a 12-bit constant.

For `lpre`, the following encoding is used, with each character representing one bit:
`0001 0iii iiii iiii  iiii iiii iiii iiii`, where
- `i` is a 27-bit constant.

For `cmpxchg` and `xchg`, the following encoding is used, with each character representing one bit:
`0001 100l bbbb aaaa`, where
- `l` encodes whether to lock (`1`) or not lock (`0`)
- `b` encodes register `rB`
- `a` encodes register `rA`

This instruction acts as `cmpxchg` if `index` is in effect, using `indexRegA` as the `expected` (or `rC`) value, but otherwise this instruction acts as `xchg`. More on `index` later in this document.

The assembly syntax of these instructions is as follows:
- For `l=0` (without lock):
  - `cmpxchg [rA], rC, rB`
  - `xchg rA, rB`
- For `l=1` (with lock):
  - `cmpxchg.l [rA], rC, rB`
  - `xchg.l [rA], rB`

`cmpxchg`/`cmpxchg.l` sets the `Z` flag to `1` upon success and `0` upon failure.

II.A.1 `pre` and `lpre`

`pre` and `lpre` are mechanisms by which immediates larger than normal can be used, essentially acting like variable width instructions.

There is no mechanism in the assembly language itself to use `pre` or `lpre` as instructions. Instead, it is expected that the assembler or linker will be the one to insert `pre` or `lpre` as needed if an immediate is too large for a particular instruction.

For non-branch instructions:
- If `pre` is used, the immediate field of the `pre` instruction will form bits `[16:5]` of the immediate of the next non-`index` instruction. The 17-bit immediate will then be sign-extended to 32 bits. Sign-extension will be performed on the 17-bit immediate even if, had there been no `pre`, the 5-bit immediate would have been zero-extended.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:5]` of the immediate of the next non-`index` instruction.

For branch instructions (group 3):
- if `pre` is used, the immediate field of the `pre` instruction will form bits `[20:9]` of the immediate of the next non-`index` instruction. The 21-bit immediate will then be sign-extended to 32 bits.
- On the other hand, if `lpre` is used, the immediate field of the `lpre` instruction will form bits `[31:9]` of the immediate of the next non-`index` instruction.

II.A.2 Handling of `pre`, `lpre`, and `index`

When a `pre` or `lpre` instruction is found, `pre` or `lpre` will be considered to be "in effect". This condition lasts for one or two instructions after the `pre` or `lpre` instruction, depending on whether or not `index` was in effect.

`index` is an instruction (defined later) that allows a subsequent load or store instruction to perform `base_reg + indexRegSum` indexing. When an `index` instruction is found, it is considered to be in effect. Like `pre` and `lpre`, `index` is inserted automatically by the assembler.

`index` can be combined with `pre` or `lpre`, and it does not matter whether `index` or the `pre`/`lpre` instruction came first.

When `pre`, `lpre`, or `index` is in effect, IRQs will not be serviced.

Pseudo code for handling the how `pre`, `lpre`, and `index` are to be handled regarding whether or not they are "in effect" is as follows:

```
function handlePreLpreIndex(input in, output out) {
  when (!out.instr.isLpre()) {
    when (in.instr.isPre()) {
      when (
        in.state.pre.haveIt
        || in.state.lpre.haveIt
      ) {
        // invalid instruction
        out.state.pre.haveIt := False
        out.state.lpre.haveIt := False
        out.state.index.haveIt := False
```

```
      out.instr := NOP
      out.state.canServiceIrq := True
    } otherwise {
      out.state.pre.haveIt := True
      out.state.pre.data := in.instr[11:0]
      out.instr := in.instr
      out.state.canServiceIrq := False
    }
  } elsewhen (in.instr.isLpre()) {
    when (
      in.state.pre.haveIt
      || in.state.lpre.haveIt
    ) {
      // invalid instruction
      out.state.pre.haveIt := False
      out.state.lpre.haveIt := False
      out.state.index.haveIt := False
      out.instr := NOP
      out.state.canServiceIrq := True
    } otherwise {
      out.state.lpre.haveIt := True
      out.instr := in.instr
      out.state.canServiceIrq := False
      out.state.lpre.data[26:16] := in.instr[10:0]
    }
  } elsewhen (in.instr.isIndex()) {
    when (
      in.state.index.haveIt
    ) {
      // invalid instruction
      out.state.pre.haveIt := False
      out.state.lpre.haveIt := False
      out.state.index.haveIt := False
      out.instr := NOP
      out.state.canServiceIrq := True
    } otherwise {
      out.state.index.haveIt := True
      out.instr := in.instr
      out.state.canServiceIrq := False
    }
  } otherwise {
    out.state.canServiceIrq := !(
      in.state.pre.haveIt
      || in.state.lpre.haveIt
      || in.state.index.haveIt
    )
    out.instr := instr

    // Whenever we see an instruction other than `pre`, `lpre`, or `index`,
    // that means those instructions stop being "in effect".
    out.state.pre.haveIt := False
    out.state.lpre.haveIt := False
    out.state.index.haveIt := False
  }
} otherwise { // when (out.instr.isLpre())
  out.state.lpre.data[15:0] := in.instr
}
}
```

II.B Instruction Group 1

The following encoding is used, with each character representing one bit:

`001i iiii oooo aaaa`, where

- i is a 5-bit sign-extended or zero-extended immediate, and is denoted `simm` when sign-extended or `imm` when zero-extended. Also, `simm` or `imm` can be expanded with `pre` or `lpre`.
- `a` encodes register `rA`
- `o` is the opcode

Here is a list of instructions from this encoding group.
- Opcode `0x0`: add rA, #simm
- Opcode `0x1`: add rA, pc, #simm
  - Effect: `rA <= pc + simm + 2;`
- Opcode `0x2`: add rA, sp, #simm
- Opcode `0x3`: add rA, fp, #simm
- Opcode `0x4`: cmp rA, #simm
  - Effect: Compare `rA` to `simm`.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x5`: cpy rA, #simm
  - Effect: Copy an immediate value into `rA`
- Opcode `0x6`: lsl rA, #imm
  - Effect: Logical shift left
- Opcode `0x7`: lsr rA, #imm
  - Effect: Logical shift right
- Opcode `0x8`: asr rA, #imm
  - Effect: Arithmetic shift right
- Opcode `0x9`: and rA, #simm
  - Effect: Bitwise AND
- Opcode `0xa`: orr rA, #simm
  - Effect: Bitwise OR
- Opcode `0xb`: xor rA, #simm
  - Effect: Bitwise XOR
- Opcode `0xc`: ze rA, #imm
  - Effect: Set `rA[31:imm]` to zero.
- Opcode `0xd`: se rA, #imm
  - Effect: Set each bit of `rA[31:imm]` to the bit `rA[imm]`.
- Opcode `0xe`: swi rA, #simm
  - Effect: Call software interrupt number `rA + simm`.
- Opcode `0xf`: swi #imm
  - Effect: Call software interrupt number `imm`.

II.C Instruction Group 2

The following encoding is used, with each character representing one bit:
`010f oooo bbbb aaaa`, where
- `o` is the opcode
- `b` encodes register `rB`
- `a` encodes register `rA`
- `f` is encoded as `0` if this instruction cannot affect flags and encoded `1` if this instruction is permitted to affect flags. Note that `cmp` is permitted to affect flags regardless of this bit.

Here is a list of instructions from this encoding group.
- Opcode `0x0`: add rA, rB
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x1`: sub rA, rB
  - Mnemonic for when flags not affected: `sub`
  - Mnemonic for when flags affected: `sub.f`

- Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x2`: `add rA, sp, rB`
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x3 add rA, fp, rB`
  - Mnemonic for when flags not affected: `add`
  - Mnemonic for when flags affected: `add.f`
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x4`: `cmp rA, rB`
  - Effect: Compare `rA` to `rB`. `cmp` is *always* able to affect flags, independent of the encoded `f` bit of the instruction.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x5`: `cpy rA, rB`
  - Mnemonic for when flags not affected: `cpy`
  - Mnemonic for when flags affected: `cpy.f`
  - Effect: Copy `rB` into `rA`
  - Affectable flags: `Z`, `N`
- Opcode `0x6`: `lsl rA, rB`
  - Mnemonic for when flags not affected: `lsl`
  - Mnemonic for when flags affected: `lsl.f`
  - Effect: Logical shift left
  - Affectable flags: `Z`, `N`
- Opcode `0x7`: `lsr rA, rB`
  - Mnemonic for when flags not affected: `lsr`
  - Mnemonic for when flags affected: `lsr.f`
  - Effect: Logical shift right
  - Affectable flags: `Z`, `N`
- Opcode `0x8`: `asr rA, rB`
  - Mnemonic for when flags not affected: `asr`
  - Mnemonic for when flags affected: `asr.f`
  - Effect: Arithmetic shift right
  - Affectable flags: `Z`, `N`
- Opcode `0x9`: `and rA, rB`
  - Mnemonic for when flags not affected: `and`
  - Mnemonic for when flags affected: `and.f`
  - Effect: Bitwise AND
  - Affectable flags: `Z`, `N`
- Opcode `0xa`: `orr rA, rB`
  - Mnemonic for when flags not affected: `orr`
  - Mnemonic for when flags affected: `orr.f`
  - Effect: Bitwise OR
  - Affectable flags: `Z`, `N`
- Opcode `0xb`: `xor rA, rB`
  - Mnemonic for when flags not affected: `xor`
  - Mnemonic for when flags affected: `xor.f`
  - Effect: Bitwise XOR
  - Affectable flags: `Z`, `N`
- Opcode `0xc`: `adc rA, rB`
  - Mnemonic for when flags not affected: `adc`
  - Mnemonic for when flags affected: `adc.f`
  - Effect: Add with Carry, using the formula `rA + rB + flags.C` to compute the value that will be written into `rA`.

- Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0xd`: `sbc rA, rB`
  - Mnemonic for when flags not affected: `sbc`
  - Mnemonic for when flags affected: `sbc.f`
  - Effect: Subtract with Borrow, using the formula `rA + (~rB) + flags.C` to compute the value that will be written into `rA`.
  - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0xe`: `cmpbc rA, rB`
  - Effect: Compare `rA` to `rB`, but with carry-in and a different effect for setting the `flags.Z`. `cmpbc` is *always* able to affect `flags`, independent of the encoded `f` bit of the instruction.
  - Note: this instruction acts much like `sbc rA, rB`, but without storing the subtraction's result into `rA`. However, this instruction sets the `Z` flag to `prev(flags.Z) AND ((rA + (~rB) + flags.C) == 0)`
  - Affectable flags: `Z`, `C`, `V`, `N`

II.D Instruction Group 3: Relative Branches

The following encoding is used, with each character representing one bit:
`011i iiii iiii oooo`, where

- `i` is a 9-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `o` is the opcode

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `bl simm`
  - Name: Branch and Link
  - Description: Relative call
  - Effect: `lr <= pc + 2; pc <= pc + simm + 2;`
- Opcode `0x1`: `bra simm`
  - Name: BRanch Always
  - Description: Unconditional relative branch
  - Effect: `pc <= pc + simm + 2;`
- Opcode `0x2`: `beq simm`
  - Name: Branch if EQual
  - Effect: `if (flags.Z) pc <= pc + simm + 2;`
- Opcode `0x3`: `bne simm`
  - Name: Branch if Not Equal
  - Effect: `if (!flags.Z) pc <= pc + simm + 2;`
- Opcode `0x4`: `bmi simm`
  - Name: Branch if MInus
  - Effect: `if (flags.N) pc <= pc + simm + 2;`
- Opcode `0x5`: `bpl simm`
  - Name: Branch if PLus
  - Effect: `if (!flags.N) pc <= pc + simm + 2;`
- Opcode `0x6`: `bvs simm`
  - Name: Branch if oVerflow Set
  - Effect: `if (flags.V) pc <= pc + simm + 2;`
- Opcode `0x7`: `bvc simm`
  - Name: Branch if oVerflow Clear
  - Effect: `if (!flags.V) pc <= pc + simm + 2;`
- Opcode `0x8`: `bgeu simm`
  - Name: Branch if Greater than or Equal Unsigned
  - Effect: `if (flags.C) pc <= pc + simm + 2;`
- Opcode `0x9`: `bltu simm`
  - Name: Branch if Less Than Unsigned

- Effect: `if (!flags.C) pc <= pc + simm + 2;`
- Opcode `0xa`: `bgtu simm`
  - Name: Branch if Greater Than Unsigned
  - Effect: `if (flags.C AND !flags.Z) pc <= pc + simm + 2;`
- Opcode `0xb`: `bleu simm`
  - Name: Branch if Less than or Equal Unsigned
  - Effect: `if (!flags.C OR flags.Z) pc <= pc + simm + 2;`
- Opcode `0xc`: `bges simm`
  - Name: Branch if Greater than or Equal Signed
  - Effect: `if (!(flags.N XOR flags.V)) pc <= pc + simm + 2;`
- Opcode `0xd`: `blts simm`
  - Name: Branch if Less Than Signed
  - Effect: `if (flags.N XOR flags.V) pc <= pc + simm + 2;`
- Opcode `0xe`: `bgts simm`
  - Name: Branch if Greater Than Signed
  - Effect:

    ```
    if (!(flags.N XOR flags.V) AND !flags.Z)
       pc <= pc + simm + 2;
    ```
- Opcode `0xf`: `bles simm`
  - Name: Branch if Less than or Equal Signed
  - Effect:

    ```
    if ((flags.N XOR flags.V) OR flags.Z)
       pc <= pc + simm + 2;
    ```

II.E Instruction Group 4

The following encoding is used, with each character representing one bit:
`1000 oooo bbbb aaaa`, where

- `o` is the opcode
- `b` encodes register `rB` or register `sC`, where `sC` is one of the special-purpose registers
- `a` encodes register `rA` or register `sA`, where `sA` is one of the special-purpose registers

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `jl rA`
  - Effect: `lr <= pc + 2; pc <= rA;`
- Opcode `0x1`: `jmp rA`
  - Effect: `pc <= rA;`
- Opcode `0x2`: `jmp ira`
  - Effect: `pc <= ira;`
- Opcode `0x3`: `reti`
  - Effect: enables IRQs (by copying `0x1` into `ie`) and performs `pc <= ira;`
- Opcode `0x4`: `ei`
  - Effect: copy `1` into `ie`.
- Opcode `0x5`: `di`
  - Effect: copy `0` into `ie`.
- Opcode `0x6`: `push rA, rB`
  - Effect: pushes `rA` onto the stack, using `rB` as the stack pointer, post-decrementing `rB`.
  - This instruction does nothing when `rA` is the same register as `rB`.
  - Note: As a pseudo instruction, omitting "`, rB`" will automatically select `sp` as the particular stack pointer.
- Opcode `0x7`: `push sA, rB`
  - Effect: pushes `sA` onto the stack, using `rB` as the stack pointer, post-decrementing `rB`.
  - Note that `sA` is considered to be 32-bit for the purpose of the store to memory and decrementing `rB`, even if `sA` is `flags` or `ie`.

- Note: As a pseudo instruction, omitting ", rB" will automatically select `sp` as the particular stack pointer.
- Opcode `0x8`: `pop rA, rB`
  - Effect: pops `rA` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
  - This instruction does nothing when `rA` is the same register as `rB`.
  - Note: As a pseudo instruction, omitting ", rB" will automatically select `sp` as the particular stack pointer.
- Opcode `0x9`: `pop sA, rB`
  - Effect: pops `sA` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
  - Note that `sA` is considered to be 32-bit for the purpose of the load from memory and incrementing `rB`, even if `sA` is `flags` or `ie`.
  - Note: As a pseudo instruction, omitting ", rB" will automatically select `sp` as the particular stack pointer.
- Opcode `0xa`: `pop pc, rB`
  - Effect: pops `pc` off the stack, using `rB` as the stack pointer, pre-incrementing `rB`.
  - Note: As a pseudo instruction, omitting ", rB" will automatically select `sp` as the particular stack pointer.
- Opcode `0xb`: `mul rA, rB`
  - Effect: `rA <= rA * rB;`
- Opcode `0xc`: `udiv rA, rB`
  - Effect: `rA <= u32(rA) / u32(rB);`
  - Note: If `index` is in effect, this instruction becomes
  - Note: This instruction has another form, `udivmod rA, rB, rC`, where `rC` is used as the modulo output register of the 32-bit divide. This form of this instruction is encoded as `index rC, r0` followed by `udiv rA, rB`
- Opcode `0xd`: `sdiv rA, rB`
  - Effect: `rA <= s32(rA) / s32(rB);`
  - Note: This instruction has another form, `sdivmod rA, rB, rC`, where `rC` is used as the modulo output register of the 32-bit divide. This form of this instruction is encoded as `index rC, r0` followed by `sdiv rA, rB`
- Opcode `0x0e`: `lumul rA, rB`
  - Effect: This instruction multiplies `rA` by `rB`, performing an unsigned 32-bit by 32-bit -> 64-bit multiply, storing result in
    - `concat{r0, r1}` (only if `index` is not in effect).
    - `concat{index_rA, index_rB}` (only if `index` is in effect), which uses the assembly syntax `lumul rC, rD, rA, rB`, where
      - `rC` becomes `index_rA`, and
      - `rD` becomes `index_rB`
- Opcode `0x0f`: `lsmul rA, rB`
  - Effect: This instruction multiplies `rA` by `rB`, performing a signed 32-bit by 32-bit -> 64-bit multiply, storing result in
    - `concat{r0, r1}` (only if `index` is not in effect).
    - `concat{index_rA, index_rB}` (only if `index` is in effect), which uses the assembly syntax `lumul rC, rD, rA, rB`, where
      - `rC` becomes `index_rA`, and
      - `rD` becomes `index_rB`
- Opcode `0x10`: `udiv64 rA, rB`
  - Effect: performs a 64-bit by 64-bit unsigned division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),

- Note: This instruction has another form, `udivmod64 rA, rB, rC, rD`, where `concat(rC, rD)` is used as the 64-bit register pair in which the modulo result of the 64-bit divide is stored. This form of this instruction is encoded as `index rC, rD` followed by `udiv64 rA, rB`
- Opcode `0x11`: `sdiv64 rA, rB`
  - Effect: performs a 64-bit by 64-bit signed division of `concat{rA, r{A + 1}}` by `concat{rB, r{B + 1}}`, storing 64-bit result in `concat{rA, r{A + 1}}`.
  - Note: This instruction operates as if `rA` were encoded with (`A[0x0] == 0b0`) and `rB` were encoded with (`B[0x0] == 0b0`),
  - Note: This instruction has another form, `sdivmod64 rA, rB, rC, rD`, where `concat(rC, rD)` is used as the 64-bit register pair in which the modulo result of the 64-bit divide is stored. This form of this instruction is encoded as `index rC, rD` followed by `sdiv64 rA, rB`
- Opcode `0x12`: `ldub rA, [rB]`
  - Effect: Load an 8-bit value from memory at the the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),

    then zero-extend the 8-bit value to 32 bits, then write the zero-extended 32-bit value into `rA`.
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `ldub rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `ldub rA, [rB, #simm]`
- Opcode `0x13`: `ldsb rA, [rB]`
  - Effect: Load an 8-bit value from memory the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),

    then sign-extend the 8-bit value to 32 bits, then write the sign-extended 32-bit value into `rA`.
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `ldsb rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `ldsb rA, [rB, #simm]`
- Opcode `0x14`: `lduh rA, [rB]`
  - Effect: Load a 16-bit value from memory at the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),

    then zero-extend the 16-bit value to 32 bits, then write the zero-extended 32-bit value into `rA`.
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `lduh rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `lduh rA, [rB, #simm]`
- Opcode `0x15`: `ldsh rA, [rB]`
  - Effect: Load a 16-bit value from memory at the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),

    then sign-extends the 16-bit value to 32 bits, then writes the sign-extended 32-bit value into `rA`.
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `ldsh rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `ldsh rA, [rB, #simm]`
- Opcode `0x16`: `ldr rA, [rB]`
  - Effect: Store `rA` to memory at the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),

- Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `ldr rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `ldr rA, [rB, #simm]`
- Opcode `0x18`: `stb rA, [rB]`
  - Effect: Store `rA[7:0]` to memory at the the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `stb rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `stb rA, [rB, #simm]`
- Opcode `0x19`: `sth rA, [rB]`
  - Effect: Store `rA[15:0]` to memory at the the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `sth rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `sth rA, [rB, #simm]`
- Opcode `0x1a`: `str rA, [rB]`
  - Effect: Store `rA` to memory at the the effective address computed as either
    - `rB` (only if `index` is **not** in effect)
    - or `<indexRegSum>` (only if `index` is in effect),
  - Shorthand for having the assembler insert an `index rB, rC` instruction before this one: `str rA, [rB, rC]`
  - Shorthand for having the assembler insert an `index rB, #simm` instruction before this one: `str rA, [rB, #simm]`
- Opcode `0x1c`: `cpy rA, sB`
  - Effect: `rA <= sB;`
- Opcode `0x1d`: `cpy sA, rB`
  - Effect: `sA <= rB;`
- Opcode `0x1e`: `cpy sA, sB`
  - Effect: `sA <= sB;`

II.F Instruction Group 5, Subgroup `0b0`: `index rA, rB`

The following encoding is used, with each character representing one bit:
`1010 000r bbbb aaaa`, where
- `r` is a reserved bit value, which should be zero until the 64-bit version of this spec is defined. defined
- `b` encodes register `rB`
- `a` encodes register `rA`

The one instruction from this encoding group is `index rA, rB`
- Effect: Performs the following:
  - `indexRegSum <= indexRaSimm.haveIt ? indexRegSum + rB : rA + rB;`
  - `indexRegA <= rA;`
  - `indexRegB <= rB;`

  and stores that `indexRaRb` is in effect.
- Note: If this prefix occurs when `indexRaSimm` is already in effect, then the overall instruction becomes a NOP, and `indexRaRb`, `indexRaSimm`, `pre`, and `lpre` will all no longer be in effect.

II.G Instruction Group 5, Subgroup `0b1`: `index rA, #simm`

The following encoding is used, with each character representing one bit:
`1011 riii iiii aaaa`, where

- `r` is a reserved bit value, which should be encoded as zero until the 64-bit version of this spec is defined.
- `i` is a 7-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`
- `a` encodes register `rA`

The one instruction from this encoding group is `index rA, #simm`
- Effect: Performs the following:
  - `indexRegSum <= indexRaRb.haveIt ? indexRegSum + simm : rA + simm;`
  - `indexRegA <= rA;`
  - `indexRegB <= simm;`

  and stores that `indexRaSimm` is in effect.
- Note: If this prefix occurs when `indexRaSimm` is already in effect, then the overall instruction becomes a NOP, and `indexRaRb`, `indexRaSimm`, `pre`, and `lpre` will all no longer be in effect.

II.H Instruction Group 7, Subgroup `0b00`: Extra 8-bit and 16-bit Ops

The following encoding is used, with each character representing one bit:
`1110 0woo bbbb aaaa`, where
- `w` is the operation width
  - When `0b0`: 8-bit operation
  - When `0b0`: 16-bit operation
- `o` is the opcode
- `b` encodes register `rB`
- `a` encodes register `rA`

Here is a list of instructions from this encoding group.
- Opcode `0x0`:
  - `w` value `0b0`: `cmpb rA, rB`
    - Effect: Compare `rA[7:0]` to `rB[7:0]`
    - Affectable flags: `Z`, `C`, `V`, `N`
  - `w` value `0b1`: `cmph rA, rB`
    - Effect: Compare `rA[15:0]` to `rB[15:0]`
    - Affectable flags: `Z`, `C`, `V`, `N`
- Opcode `0x1`:
  - `w` value `0b0`: `lsrb rA, rB`
    - Effect: Logical shift right `rA[7:0]` by `rB`
  - `w` value `0b1`: `lsrh rA, rB`
    - Effect: Logical shift right `rA[15:0]` by `rB`
- Opcode `0x2`:
  - `w` value `0b0`: `asrb rA, rB`
    - Effect: Arithmetic shift right `rA[7:0]` by `rB`
  - `w` value `0b1`: `asrh rA, rB`
    - Effect: Arithmetic shift right `rA[15:0]` by `rB`

II.I Instruction Group 7, Subgroup `0b010`: Extra load/store instructions

The following encoding is used, with each character representing one bit:
`1110 10oo bbbb aaaa`, where
- `o` is the opcode
- `b` encodes register `rB` or `sB`
- `a` encodes register `sA`

Here is a list of instructions from this encoding group.
- Opcode `0x0`: `ldr sA, [rB]`
  - Effect: 32-bit load of `sA` from memory at address held in `rB`.
- Opcode `0x1`: `ldr sA, [sB]`

- Effect: 32-bit load of `sA` from memory at address held in `sB`.
- Opcode `0x2`: str sA, [rB]
  - Effect: 32-bit store of `sA` to memory at address held in `rB`.
- Opcode `0x3`: str sA, [sB]
  - Effect: 32-bit store of `sA` to memory at address held in `sB`.

II.J Instruction Group 7, Subgroup `0b0110`: Icache Reload Instruction

The following encoding is used, with each character representing one bit:

`1110 110i iiii aaaa`, where

- `i` is a 5-bit sign-extended immediate, which can be expanded by `pre` or `lpre`, and is denoted `simm`.
- `a` encodes register `rA`

The one instruction from this encoding group is `icreload [rA, #simm]`. This instruction forcibly reloads the icache line at an effective address computed as `rA + simm` (only if `index` is **not** in effect) or `<indexRegSum> + simm` (only if `index` is in effect), using the sign-extended form of `simm`. As a side note, this instruction will attempt to read the data from dcache, and will only read from memory if dcache doesn't have that data. Icache lines and dcache lines are the same length, 64 bytes.

The following pseudo instructions exist within the assembler for this encoding group:

- `icreload [rA]`
  - This pseudo instruction assembles to the following:
    - `icreload [rA, #0x0]`
- `icreload [rA, rC]`
  - This pseudo instruction assembles to the following:
    - `index rA, rC`
    - `icreload [r0, #0x0]`
- `icreload [rA, rC, #simm]`
  - This pseudo instruction assembles to the following:
    - `index rA, rC`
    - `icreload [r0, #simm]`

II.K Instruction Group 7, Subgroup `0b01110`: Icache Flush Instruction

The following encoding is used, with each character representing one bit:

`1110 1110 0000 0000`

The one instruction from this encoding group is `icflush`, which invalidates the instruction cache entirely.